

# PantaRay:

## Fast Ray-traced Occlusion Caching of Massive Scenes

J. Pantaleoni, L. Fascione, M. Hill, T. Aila

Marie-Lena Eckert



# Agenda

---

- Introduction
  - Motivation
  - Basics
- PantaRay
  - Accelerating structure generation
  - Massively parallel ray tracer
- Conclusion
  - Summary
  - Critics

# Introduction: Motivation

---

- Computer-generated imagery (CGI), rendering:

3D model



2D image



# Introduction: Motivation

---

- Avatar: unprecedented complexity






# Introduction: Motivation

---

- Avatar: level of detail (LOD)




# Introduction: Basics

---

- Rendering equation


$$L_o(x, \omega_o) = \int_{\Omega^+} L_i(x, \omega) \rho(x, \omega, \omega_o) V(x, \omega) \langle \omega, \hat{n} \rangle d\omega$$



# Introduction: Basics

- Rendering equation


$$L_o(x, \omega_o) = \int_{\Omega^+} L_i(x, \omega) \rho(x, \omega, \omega_o) V(x, \omega) \langle \omega, \hat{n} \rangle d\omega$$



# Introduction: Basics

---

- Image-based lighting (IBL) for global illumination





# Introduction: Basics

---


- Image-based lighting (IBL) for global illumination



# Introduction: Basics

- Rendering equation


$$L_o(x, \omega_o) = \int_{\Omega^+} L_i(x, \omega) \overset{\text{BRDF}}{\rho(x, \omega, \omega_o)} V(x, \omega) \langle \omega, \hat{n} \rangle d\omega$$



# Introduction: Basics

- Rendering equation

$$L_o(x, \omega_o) = \int_{\Omega^+} L_i(x, \omega) \rho(x, \omega, \omega_o) V(x, \omega) \langle \omega, \hat{n} \rangle d\omega$$




# Introduction: Basics

---

- Rendering equation

$$L_o(x, \omega_o) = \int_{\Omega^+} L_i(x, \omega) \rho(x, \omega, \omega_o) V(x, \omega) \langle \omega, \hat{n} \rangle d\omega$$






# Introduction: Basics

- Rendering equation


$$L_o(x, \omega_o) = \int_{\Omega^+} L_i(x, \omega) \rho(x, \omega, \omega_o) V(x, \omega) \langle \omega, \hat{n} \rangle d\omega$$



# Introduction: Basics

- Rendering equation



$$L_o(x, \omega_o) = \int_{\Omega^+} L_i(x, \omega) \rho(x, \omega, \omega_o) V(x, \omega) \langle \omega, \hat{n} \rangle d\omega$$



# Introduction: Basics

---


- Spherical Harmonics (SH)
  - [Ram01]: 9 coefficients, 1% error
  - Store visibility term and BRDF
  - Rendering time: only scalar product



# Introduction: Basics

- Rendering equation

$$L_o(x, \omega_o) = \rho \cdot \langle SH(L_i), SH(V\langle \omega, \hat{n} \rangle) \rangle$$





# PantaRay


---

→ Precompute and cache visibility term

- Scenes:
  - Unprecedented complexity
  - Bigger than memory
  - Varying density
- Production pipeline of Weta Digital
- Panta rei: everything flows

# PantaRay: Pipeline computation passes

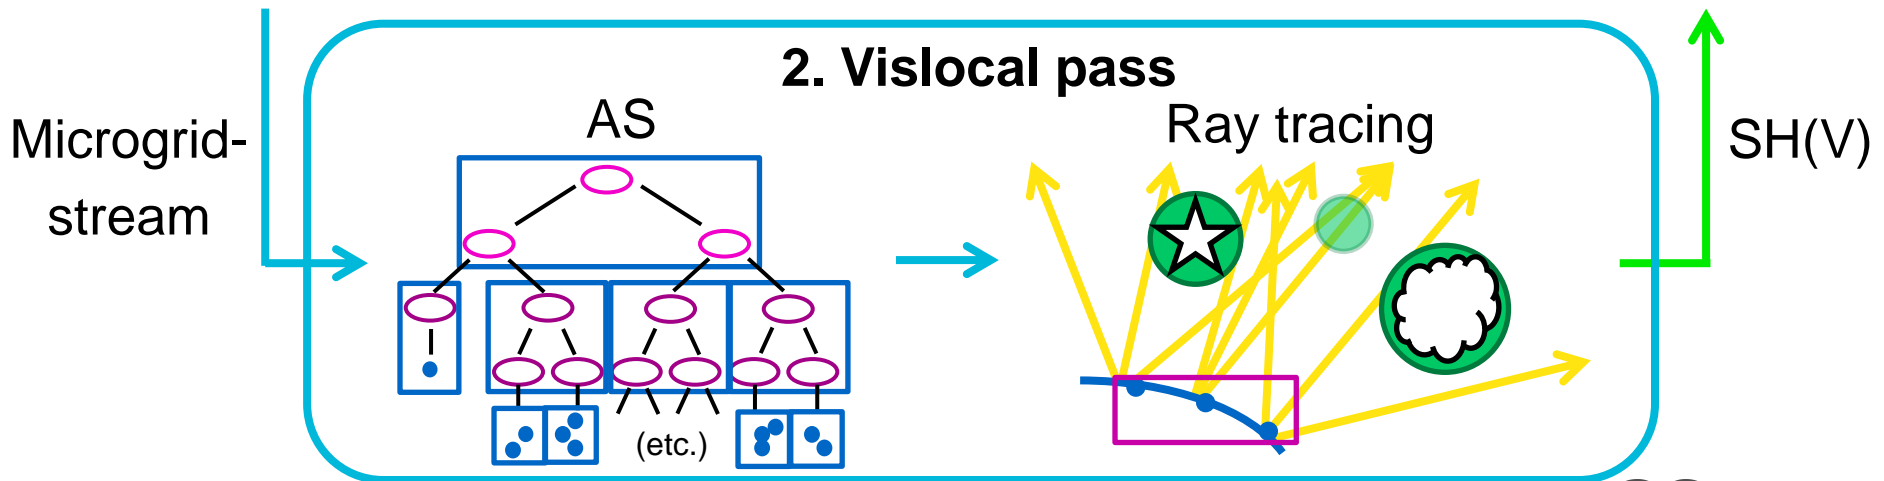
## 1. Preparation



## 3. Beauty pass

$$L_o(x, \omega_o) = \rho \cdot \langle SH(L_i), SH(\mathbf{V}\langle \omega, \hat{n} \rangle) \rangle$$

## 2. Vislocal pass



# PantaRay: Preparation

---

Microgrid stream:

Final render of the scene



Render of bake sets



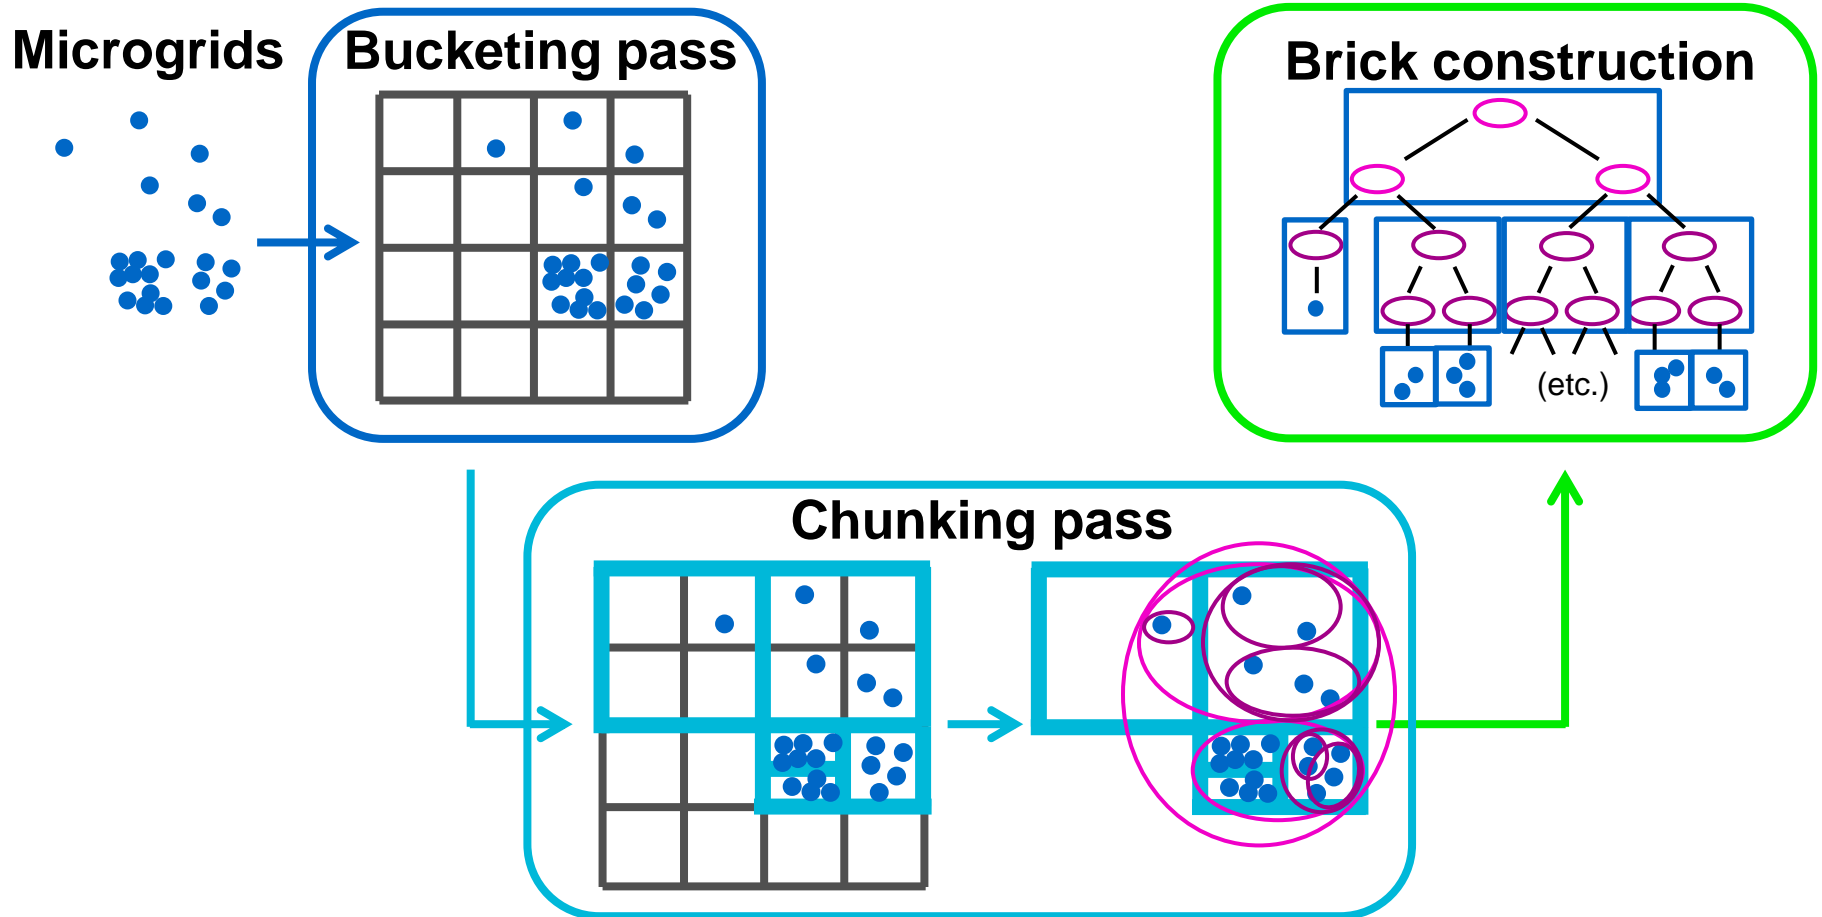
# PantaRay: Acceleration Structure

---

- Vislocal pass
  - Build Bounding Volume Hierarchy (BVH)
  - Traverse for intersection computation
    - LOD
    - Ray tracing
- Main bottleneck: I/O
  - Touch objects multiple times
  - Ten of thousands concurrent processes



# PantaRay: Acceleration Structure



# PantaRay: Ray tracing


---

- Shading of bake sets
  - Point stream into batches
  - Simple ray tracer: CPU
    - Simple reflection occlusion
    - Area light shadow shaders
  - Massively parallel ray tracer: GPU
    - Complex spherical sampling queries
      - SH occlusion
      - Indirect lighting shaders
    - Billions of points per scene → compute-intensive, LOD

# PantaRay: Ray tracing - LOD


---

- Determine and stream required bricks



# PantaRay: Ray tracing - LOD


- 3 rays emanating from each surface point
- Each ray represents solid angle  $\sigma = \frac{2\pi}{n}$ ,  $n = 3$





# PantaRay: Ray tracing - LOD


- Bricks loaded if  $\text{size}(\text{BV}(\text{brick})) \geq \text{avg}(\text{dist}(\text{rays}))$



# PantaRay: Ray tracing - LOD

---


- Partially transparent Bounding Volumes (BV)



# PantaRay: Ray tracing - LOD

---

- Custom far-field for batch 1



# PantaRay: Generate and trace rays

---

- $m$  GPUs
  - $\frac{1}{m}$  points, subset  $\rightarrow$  focus on coherence
- Trace through brick hierarchy
  - [Ail09] while-while traversal kernel
- Intersection computation
  - high scene complexity, high tree depth  $\rightarrow$  low Single Instruction Multiple Thread (SIMT) utilization
  - increase from 20% to 50-60%

# Conclusion: Summary

---

- Raised 2 orders of magnitude in terms of both speed and scene size
- Comparison AS generation
  - [Wald05]: 350M triangles: 1d
  - PantaRay: 575M micropolygons: 54min
- Custom far-field → low I/O (few MBs per batch)
- Ray tracing speed: up to 22M shaded rays per second

# Conclusion: Summary

---

I/O: 2MB per batch, trace time: 16 h, 41 m,  
8.7M shaded rays/s





# Conclusion: Critics

---

- Loss of details in computing visibility term (SH, LOD)
- Float precision: missing some intersections
- Parallel construction of AS

---

# Thank you!

# PantaRay: References

---

[Ail09]

AILA, T., AND LAINE, S. 2009. Understanding the efficiency of ray traversal on GPUs. In Proc. High-Performance Graphics, 145–149.

[Ram01]

Ravi Ramamoorthi and Pat Hanrahan. An efficient representation for irradiance environment maps. In Proceedings of the 28th annual conference on Computer graphics and interactive techniques, SIGGRAPH '01, pages 497-500, New York, NY, USA, 2001. ACM.

[Wald05]

Ingo Wald, Andreas Dietrich, and Philipp Slusallek. An interactive out-of-core rendering framework for visualizing massively complex models. In ACM SIGGRAPH 2005 Courses, SIGGRAPH '05, New York, NY, USA, 2005. ACM.

# Trace rays algorithm

---

```
trace()
1  while any active ray
2    if any active ray is in leaf node
3      perform wide_leaf_intersection() // Figure 8
4      pop traversal stack
5
6    for  $i = 0 \dots 8$  // short while loop
7      if node is a new brick
8        if new brick is not loaded
9          report intersection with new brick's bbox
10         pop traversal stack
11       else
12         jump to new brick's root
13     else if node is not leaf
14       traverse to next node
15     else
16       break // node is a leaf
```

# Compute intersections algorithm

```
wide_leaf_intersection()
1  tidx = threadIdx.x; // [0,32)
2  // count the number of ray-primitive tasks.
3  // haveLeaf indicates if the current node is a leaf
4  [lo,hi) = scan(haveLeaf ? numPrims : 0);
5  numIssect = hi from thread 31 // how many left?
6
7  while (numIssect > 0)
8  {
9      // select up to 32 ray-primitive tasks
10     foreach primitive p, with lo+p ∈ [0,32)
11         write (tidx,p) pair to shared[lo+p]
12
13     // get a ray-primitive task to execute
14     (srcThread,p2) = shared[tidx];
15
16     // copy needed variables from "srcThread" thread
17     foreach 32bit variable var needed in intersection
18         shared[tidx] = var; // write my own variable
19         var2 = shared[srcThread]; // read from "srcThread"
20
21     // intersection using vars copied from "srcThread"
22     if (tidx < numIssect)
23         shared[tidx] = intersectRayPrim(p2);
24
25     // collect intersections of my ray
26     foreach prim p, with lo+p ∈ [0,32)
27         modify ray according to shared[lo+p]
28
29     lo = lo-32; hi = hi-32; numIssect = numIssect-32;
30 }
```